

# Design of Experiments

Arnaud Legrand

Performance Evaluation Lecture  
UFRGS, Porto Alegre, August 2015

# Outline

## ① Design of Experiments

## ② Factorial studies

- 2-level Factorial Studies

- ANOVA

- Fractional design and Screening

- General factorial designs

## ③ Model Investigation

- Designs

- Exploiting and Reducing Variance

- Discussing the Shape of the Model

## ④ Model Estimation

- Optimal Designs

## ⑤ Conclusion

# Key concepts

There are two key concepts:

replication and randomization

You replicate to increase reliability. You randomize to reduce bias.

**If you replicate thoroughly and randomize properly,  
you will not go far wrong.**

# Key concepts

There are two key concepts:

**replication** and **randomization**

You replicate to **increase reliability**. You randomize to **reduce bias**.

**If you replicate thoroughly and randomize properly,  
you will not go far wrong.**

*It doesn't matter if you cannot do your own advanced statistical analysis. If you designed your experiments properly, you may be able to find somebody to help you with the statistics.*

*If your experiments is not properly designed, then no matter how good you are at statistics, you experimental effort will have been wasted.*

**No amount of high-powered statistical analysis can turn a bad experiment into a good one.**

Other important concepts:

- **Pseudo-replication**
- **Experimental** vs. **observational** data

# Select the problem to study

Clearly define the kind of **system** to study, the kind of **phenomenon** to observe (state, evolution of state through time), the kind of **study** to conduct (descriptive, exploratory, prediction, hypothesis testing, ...)

This is quite important as the set of experiments to perform will be completely different when:


- studying the stabilization of a peer-to-peer algorithm under a high churn
- trying to compare various scheduling algorithms or code versions
- modeling the response time of a server under a workload close to the server saturation
- ...

This first step enables to decide on **which kind of design** should be used

# Define the set of relevant *response*

The system under study is generally modeled though a **black-box** model:

- some **output** variable/**response**( $y$ )
- some inputs are fully unknown
- some **input variables** ( $x_1, \dots, x_p$ ) are **controllable**
- whereas some others ( $z_1, \dots, z_q$ ) are **uncontrollable**



In our case, the response could be:


- the makespan of a scheduling algorithm
- the amount of messages exchanged in a peer-to-peer system
- the convergence time of distributed algorithm
- the average length of a random walk
- the amount of energy or of memory used

Some of these metrics are the result of complex aggregation of measurements so they should be **carefully recorded** to check their correctness

# Determine the set of relevant *factors* or *variables*

The system under study is generally modeled though a **black-box** model:

- some **output** variable/**response**( $y$ )
- some inputs are fully unknown
- some **input variables** ( $x_1, \dots, x_p$ ) are **controllable**
- whereas some others ( $z_1, \dots, z_q$ ) are **uncontrollable**



Typical controllable variables could be:

- the heuristic used (e.g., FIFO, HEFT, ...)
- one of their parameters (e.g., replication factor, a threshold, ...)
- the size of the platform
- the degree of heterogeneity
- the version of the compiler

Uncontrollable variables could be:

- temperature, humidity, moon phase, road surface conditions
- someone using the machine and interfering with the experiment

You should **carefully record** all the factors you can think of

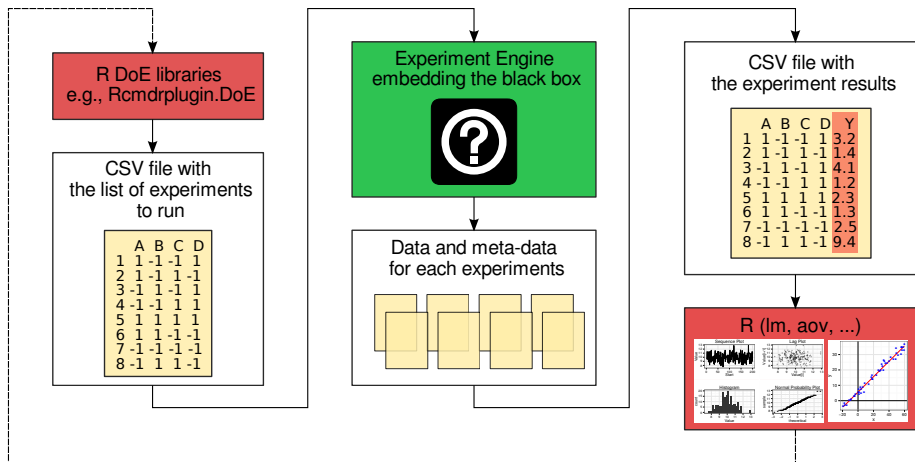
# Typical case studies

The typical case studies defined in the first step could include:

- Determining which variables are most influential on the response  $y$  (factorial designs, screening designs, analysis of variance)
  - Allows to distinguish between primary factors whose influence on the response should be modeled and secondary factors whose impact should be averaged
  - Allows to determine whether some factors interact in the response
- Devise an analytical model of the response  $y$  as a function of the primary factors  $x$  (regression, lhs designs)
- Fit a an analytical model (regression, response surface methodology, optimal designs)
  - Can then be used to determine where to set the primary factors  $x$  so that response  $y$  is always close to a desired value or is minimized/maximized
- Determining where to set the primary factors  $x$  so that variability in response  $y$  is small i.e., so that the effect of uncontrollable variables  $z_1, \dots, z_q$  is minimized (robust designs, Taguchi designs)



# General Workflow



# Outline

## ① Design of Experiments

## ② Factorial studies

- 2-level Factorial Studies

- ANOVA

- Fractional design and Screening

- General factorial designs

## ③ Model Investigation

- Designs

- Exploiting and Reducing Variance

- Discussing the Shape of the Model

## ④ Model Estimation


- Optimal Designs

## ⑤ Conclusion

# Linear Regression

$$Y = a + bX + \varepsilon$$

- $Y$  is the **response variable**
- $X$  is a continuous **explanatory variable**
- $a$  is the **intercept**
- $b$  is the **slope**
- $\varepsilon$  is some **noise**



When there are 2 explanatory variables:



$$Y = a + b^{(1)}X^{(1)} + b^{(2)}X^{(2)} + b^{(1,2)}X^{(1)}X^{(2)} + \varepsilon$$

$\varepsilon$  is generally assumed to be independent of  $X^{(k)}$ , hence it needs to be **checked** once the regression is done

- Although your phenomenon is not linear, the linear model helps for **initial investigations** (as a first crude approximation)
- You should always wonder whether there is a way of looking at your problem where it is linear

## 2-level factorial designs

- 1 Decide a **low** and a **high** value for

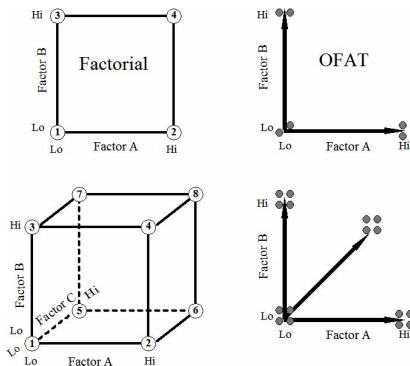


The different values are by convention encoded with  $-1$  and  $1$  but these are **not real numbers**

- 2 Test **every** ( $2^P$ ) **combination** of high and low values, possibly replicating for each combination.

By varying everything, we can detect **interactions** right away

# The downsides of the *One Factor At a Time* approach



- Only a very small fraction of the space is covered (bias) 😞
- Interaction between factors cannot be estimated 😞
- Each replication allows to improve the estimation quality of only one factor, hence it requires more runs to have good estimates of all factors 😞

Unless dealing with a very simple problem, it is always better to **change parameters all together** than change parameters **One Factor at a Time**

# Generating a $2^p$ Design

```
1 library(FrF2)
2 d1 = FrF2(nruns=8 ,nfactors=3 , blocks=1 , replications = 2,
3 randomize= TRUE, seed= 26052 ,
4 factor.names=list(A=c(-1,1), B=c(-1,1), C=c(-1,1))); d1 ;
```

```
1 creating full factorial with 8 runs ...
```

```
2
3 run.no run.no.std.rp A B C
4 1 1 2.1 1 -1 -1
5 2 2 6.1 1 -1 1
6 3 3 3.1 -1 1 -1
7 4 4 5.1 -1 -1 1
```

```
8 ...
9 15 15 1.2 -1 -1 -1
10 16 16 4.2 1 1 -1
```

```
11 class=design, type= full factorial
```

```
12 NOTE: columns run.no and run.no.std.rp are annotation, not part of t
```

How can we analyze something like this?

# Outline

## ① Design of Experiments

## ② Factorial studies

2-level Factorial Studies

ANOVA

Fractional design and Screening

General factorial designs

## ③ Model Investigation

Designs

Exploiting and Reducing Variance

Discussing the Shape of the Model

## ④ Model Estimation

Optimal Designs

## ⑤ Conclusion

# Confidence

If we had only 1 factor with 2 levels ( $2^1$  design), the analysis would simply amount to **compute confidence intervals** or more precisely to **test whether  $\mu_{A=Low} = \mu_{A=High}$**  or not (t-test)

(if few observations are available we would have to make the C.I wider and use the Student distribution)

But when having more factors and/or levels, we want to test whether **some of the combinations** have a significantly different expected value

Number of comparisons	2	3	4	5	6
Nominal Type I error	5%	5%	5%	5%	5%
Actual overall Type I error	5%	12.2%	20.3%	28.6%	36.6%

(See 16.1.5 of *Practical Regression and Anova using R* by Julian Faraway)



# Quick illustration of the difficulty of multiple testing


```
1 se = .1; mean = 7 ;
2 N = 10000;
3 df = data.frame(
4   x1=rnorm(N,mean=mean,sd=se),
5   x2=rnorm(N,mean=mean,sd=se),
6   x3=rnorm(N,mean=mean,sd=se))
7 df$eq1 = abs(df$x1-mean)<2*se
8 df$eq2 = abs(df$x1-df$x2)<2*se
9 df$eq3 = abs(df$x1-df$x2)<2*se &
10   abs(df$x1-df$x3)<2*se &
11   abs(df$x2-df$x3)<2*se
12
13 mean(df$eq1)
14 mean(df$eq2)
15 mean(df$eq3)
```

```
1 [1] 0.9538
2 [1] 0.8435
3 [1] 0.6596
```

# Analysis of Variance (ANOVA)

ANOVA (ANalysis Of VAriance) enable to discriminate real effects from noise

- Enables to prove that some parameters have little influence and can be randomized over (possibly with a more elaborate model)
- Decomposes variance:



- Assumes identical standard deviation for the populations (homoscedastic)
- Multiple tests at once (assuming normality):  $\mu_{A=Low,*} - \mu_{A=High,*} = 0$ ,

$$\mu_{B=Low,*} - \mu_{B=High,*} = 0, \dots$$

# ANOVA and F-statistic

The ANOVA produces an **F-statistic**, the ratio of the **variance calculated among the means** to the **variance within the samples**.

- If the group means are drawn from populations with the same mean values, the **variance between the group means** should be **lower** than the **variance of the samples**
- A higher ratio therefore implies that the samples were drawn from populations with different mean values

# ANOVA and F-statistic

The ANOVA produces an **F-statistic**, the ratio of the **variance calculated among the means** to the **variance within the samples**.

- If the group means are drawn from populations with the same mean values, the **variance between the group means** should be **lower** than the **variance of the samples**
- A higher ratio therefore implies that the samples were drawn from populations with different mean values

Let's work out a simple made-up example

```
1 Response = 10 + 2*as.numeric(d1$A) +  
2     3*as.numeric(d1$B)*as.numeric(d1$C) + rnorm(nrow(d1))  
3 d1 <- add.response(d1, Response, replace=TRUE)
```

I had to use `as.numeric` to interpret the `-1` and `1` as numbers whereas they were created as **factors**

# A simple ANOVA in R

```
1 d1_aov <- aov(Response ~ (A + B + C)^2, data=d1)
2 summary(d1_aov) # summary will call summary.aov
```

```
1           Df Sum Sq Mean Sq F value    Pr(>F)
2 A           1  22.98    22.98   38.318 0.000161 ***
3 B           1  68.02    68.02  113.417 2.11e-06 ***
4 C           1  77.60    77.60  129.402 1.21e-06 ***
5 A:B          1   0.44     0.44    0.728 0.415721
6 A:C          1   0.93     0.93    1.555 0.243804
7 B:C          1  14.62    14.62   24.374 0.000806 ***
8 Residuals    9   5.40     0.60
9 ---
10 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

So, all factors are significant and there is a significant interaction between B and C

# Can't I just read my linear regression as usual?

```
1 summary.lm(d1_aov)

1 Call:
2 lm.default(formula = Response ~ (A + B + C)^2, data = d1)
3
4 Residuals:
5      Min       1Q   Median       3Q      Max
6 -1.01845 -0.48073 -0.01537  0.45886  0.98771
7
8 Coefficients:
9             Estimate Std. Error t value Pr(>|t|)
10 (Intercept)   19.5912     0.1936 101.194 4.56e-15 ***
11 A1             1.1984     0.1936   6.190 0.000161 ***
12 B1             2.0618     0.1936  10.650 2.11e-06 ***
13 C1             2.2023     0.1936  11.375 1.21e-06 ***
14 A1:B1          0.1652     0.1936   0.853 0.415721
15 A1:C1          0.2415     0.1936   1.247 0.243804
16 B1:C1          0.9558     0.1936   4.937 0.000806 ***
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
19
20 Residual standard error: 0.7744 on 9 degrees of freedom
21 Multiple R-squared:  0.9716, Adjusted R-squared:  0.9527
22 F-statistic: 51.3 on 6 and 9 DF, p-value: 1.873e-06
```

# Can't I just read my linear regression as usual?

```
1 summary.lm(d1_aov)
```

```
1 Call:
2 lm.default(formula = Response ~ (A + B + C)^2, data = d1)
```

Wait, why is the formula so different?

$$10 + 2A + 3BC$$

```
9           Estimate Std. Error t value Pr(>|t|)
10 (Intercept)  19.5912      0.1936 101.194 4.56e-15 ***
11 A1           1.1984      0.1936   6.190 0.000161 ***
12 B1           2.0618      0.1936  10.650 2.11e-06 ***
13 C1           2.2023      0.1936  11.375 1.21e-06 ***
14 A1:B1        0.1652      0.1936   0.853 0.415721
15 A1:C1        0.2415      0.1936   1.247 0.243804
16 B1:C1        0.9558      0.1936   4.937 0.000806 ***
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
19
20 Residual standard error: 0.7744 on 9 degrees of freedom
21 Multiple R-squared:  0.9716, Adjusted R-squared:  0.9527
22 F-statistic: 51.3 on 6 and 9 DF, p-value: 1.873e-06
```

# Can't I just read my linear regression as usual?

```
1 summary.lm(d1_aov)
```

```
1 Call:
2 lm.default(formula = Response ~ (A + B + C)^2, data = d1)
```

Wait, why is the formula so different?

$$10 + 2A + 3BC$$

```
9           Estimate Std. Error t value Pr(>|t|)
10 (Intercept)  19.5912     0.1936 101.194 4.56e-15 ***
11 A1           1.1984     0.1936   6.190 0.000161 ***
12 B1           2.0618     0.1936  10.650 2.11e-06 ***
13 C1           2.2023     0.1936  11.375 1.21e-06 ***
14 A1:B1        0.1652     0.1936   0.853 0.415721
15 A1:C1        0.2415     0.1936   1.247 0.243804
16 B1:C1        0.9558     0.1936   4.937 0.000806 ***
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Because it treated the factors "-1" and "1" as 0 and 1...



Then how do I get the formula I expected? (1/2)

```
1 d1_lm <- lm(Response ~ (as.numeric(A) + as.numeric(B) +  
2   as.numeric(C))^2, data=d1)  
3 summary.aov(d1_lm)
```

```
1               Df Sum Sq Mean Sq F value    Pr(>F)  
2 as.numeric(A)    1  22.98   22.98   38.318 0.000161 ***  
3 as.numeric(B)    1  68.02   68.02  113.417 2.11e-06 ***  
4 as.numeric(C)    1  77.60   77.60  129.402 1.21e-06 ***  
5 as.numeric(A):as.numeric(B) 1    0.44    0.44    0.728 0.415721  
6 as.numeric(A):as.numeric(C) 1    0.93    0.93    1.555 0.243804  
7 as.numeric(B):as.numeric(C) 1   14.62   14.62   24.374 0.000806 ***  
8 Residuals       9    5.40    0.60  
9 ---  
10 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Sweet, it's the same as the previous ANOVA

## Then how do I get the formula I expected? (2/2)

```
1 summary(d1_lm) # summary will call summary.lm
1 Call:
2 lm.default(formula = Response ~ (as.numeric(A) + as.numeric(B) +
3   as.numeric(C))^2, data = d1)
4
5 Residuals:
6      Min       1Q   Median       3Q      Max
7 -1.01845 -0.48073 -0.01537  0.45886  0.98771
8
9 Coefficients:
10              Estimate Std. Error t value Pr(>|t|)
11 (Intercept)      15.4654     3.1870   4.853 0.000905 ***
12 as.numeric(A)     -0.0429     1.6878  -0.025 0.980277
13 as.numeric(B)     -2.6022     1.6878  -1.542 0.157516
14 as.numeric(C)     -2.7789     1.6878  -1.647 0.134064
15 as.numeric(A):as.numeric(B)  0.6606     0.7744   0.853 0.415721
16 as.numeric(A):as.numeric(C)  0.9658     0.7744   1.247 0.243804
17 as.numeric(B):as.numeric(C)  3.8232     0.7744   4.937 0.000806 ***
18 ---
19 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
20
21 Residual standard error: 0.7744 on 9 degrees of freedom
22 Multiple R-squared:  0.9716, Adjusted R-squared:  0.9527
23 F-statistic: 51.3 on 6 and 9 DF, p-value: 1.873e-06
```

## Then how do I get the formula I expected? (2/2)

```
1 summary(d1_lm) # summary will call summary.lm
```

Variability is too large to obtain good estimates of the true coefficients

$$10 + 2A + 3BC$$

One should anyway use other kind of designs to estimate continuous model parameters

```
10                                     Estimate Std. Error t value Pr(>|t|)
11 (Intercept)                        15.4654      3.1870   4.853 0.000905 ***
12 as.numeric(A)                      -0.0429      1.6878  -0.025 0.980277
13 as.numeric(B)                      -2.6022      1.6878  -1.542 0.157516
14 as.numeric(C)                      -2.7789      1.6878  -1.647 0.134064
15 as.numeric(A):as.numeric(B)         0.6606      0.7744   0.853 0.415721
16 as.numeric(A):as.numeric(C)         0.9658      0.7744   1.247 0.243804
17 as.numeric(B):as.numeric(C)         3.8232      0.7744   4.937 0.000806 ***
```


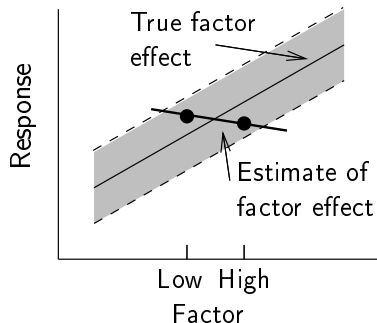
```
18 ---
19 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
20 Residual standard error: 0.7744 on 9 degrees of freedom
```

```
21 Multiple R-squared:  0.9716, Adjusted R-squared:  0.9527
```

```
22 F-statistic: 51.3 on 6 and 9 DF, p-value: 1.873e-06
```

# The difference between ANOVA and Linear Regression (3/3)




- The coding numbers are completely meaningless and influence the estimates of the slope
  - If your input parameters are numerical, go for *extreme values*, hoping the intermediate behavior is not too complicated and *consider them as factors*
- Real question: is there a *significant increase* when changing factors?
- Remember: you should use *ANOVA* for *factorial designs*, not LM
  - So don't use `summary.lm` in such cases; use `summary.aov`

## And graphically ?

```
1 MEPlot(d1, abbrev=4, select=c(1,2,3), response="Response")
```

**Main effects plot for Response**

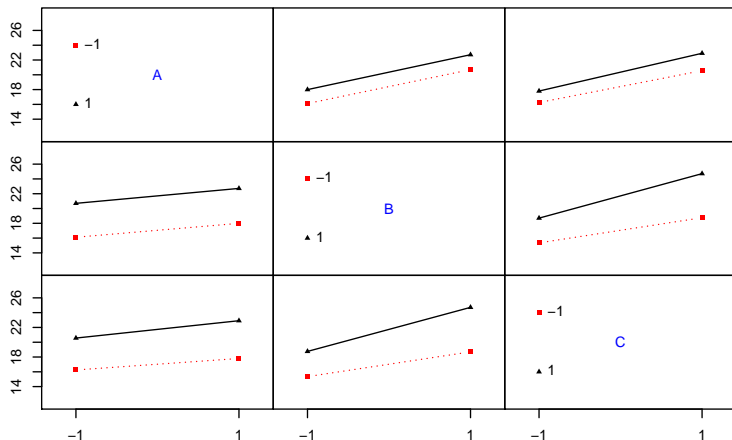


No CI on this one but we've seen that computing CIs is not straightforward  
→ rely on the `summary.aov`

# What about interactions ?

```
1 IAPlot(d1, abbrev=4, show.alias=FALSE, select=c(1,2,3))
```





**Interaction plot matrix for Response**



Again, no CI  $\leadsto$  rely on the `summary.aov`

# Checking hypothesis

```
1 layout(matrix(c(1,2,3,4),2,2)) # optional layout  
2 plot(aov(Response ~ (A + B + C)^2, data=d1))
```

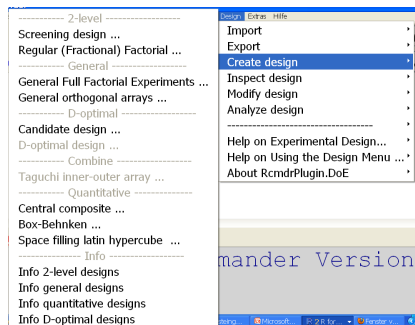


# How do you expect me to ever remember all this ?

For the R commands, there is a trick: 😊

## Use Rcmdr and Rcmdrplugin.DoE (by Ulrike Grömping)

Simply `library(Rcmdr)`...



You should only remember the principles and try to understand the underlying hypothesis

- ANOVA enables to discriminate real effects from noise in factorial experiments. *It relies on homoscedasticity and normality (or requires large number of samples)*
- 2-level factorial designs are a simple way to go and are more efficient than OFAT experiments
- Replicate thoroughly and randomize properly: you will not go far wrong



# Outline

## ① Design of Experiments

## ② Factorial studies

2-level Factorial Studies

ANOVA

Fractional design and Screening

General factorial designs

## ③ Model Investigation

Designs

Exploiting and Reducing Variance

Discussing the Shape of the Model

## ④ Model Estimation

Optimal Designs

## ⑤ Conclusion

# What if my number of factors is large ?

If  $p = 8$ , and the global variability is large, we may have to do  $r = 5$  replications, hence  $2^p \cdot r = 256 \times 5 = 1280$  experiments!!!

- Then, you need something intermediate between OFAT and a full factorial  $2^p$  design.
- It probably does not really make sense to study the joint effect of changing A, B, C, D, E, F, G, and H at the same time...

You should then go for a fractional  $2^{p-k}$  design that will still make sure the combinations are well spread and the design is well balanced

# Fractional designs

```
1 d2 = FrF2(nruns=8 ,nfactors=4 , blocks=1 , replications = 2,  
2 randomize= TRUE, seed= 26052 ,  
3 factor.names=list(A=c(-1,1), B=c(-1,1), C=c(-1,1), D=c(-1,1)));  
4 d2 ;
```

```
1      run.no run.no.std.rp  A  B  C  D  
2  1          1          2.1  1 -1 -1  1  
3  2          2          6.1  1 -1  1 -1  
4  3          3          3.1 -1  1 -1  1  
5  4          4          5.1 -1 -1  1  1  
6  5          5          8.1  1  1  1  1  
7  ...  
8 13          13          2.2  1 -1 -1  1  
9 14          14          5.2 -1 -1  1  1  
10 15          15          1.2 -1 -1 -1 -1  
11 16          16          4.2  1  1 -1 -1
```

```
12 class=design, type= FrF2
```

```
13 NOTE: columns run.no and run.no.std.rp are annotation,  
14       not part of the data frame
```

Not much gain here... Fractional designs have constraints but allow you to control how much you loose

# Saving a lot of time/money: Plackett-Burman screening designs

```
1 d3 <- pb(nruns= 20 ,n12.taguchi= FALSE ,nfactors= 20 -1, ncenter= 0 ,
2   replications= 1 ,repeat.only= FALSE ,randomize= TRUE ,seed= 26654 ,
3   factor.names=list( A=c(-1,1),B=c(-1,1),C=c(-1,1),D=c(-1,1),
4   E=c(-1,1),F=c(-1,1),G=c(-1, 1),H=c(-1,1),J=c(-1,1),K=c(-1,1),
5   L=c(-1,1),M=c(-1,1),N=c(-1,1),O=c(-1,1),P=c(-1,1) ) ) ; d3
```

	A	B	C	D	E	F	G	H	J	K	L	M	N	O	P	e1	e2	e3	e4
1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	-1	1	1	-1	-1	1
2	-1	1	-1	-1	-1	-1	1	1	-1	1	1	-1	-1	1	1	1	1	-1	1
3	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
4	1	-1	-1	1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	-1	1
5	-1	-1	1	1	-1	1	1	-1	-1	1	1	1	1	-1	1	-1	1	-1	-1
6	-1	-1	1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	-1	1	1
7	1	1	-1	1	1	-1	-1	1	1	1	1	-1	1	-1	1	-1	-1	-1	-1
8	-1	1	1	-1	1	1	-1	-1	1	1	1	1	-1	1	-1	1	-1	-1	-1
9	1	-1	1	-1	-1	-1	-1	1	1	-1	1	1	-1	-1	1	1	1	1	-1
10	1	1	-1	-1	1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	-1
11	-1	1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	-1	1	1	-1
12	1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	-1	1	1	-1	-1
13	-1	-1	-1	1	1	-1	1	1	-1	-1	1	1	1	1	-1	1	-1	1	-1
14	1	-1	-1	-1	-1	1	1	-1	1	1	-1	-1	1	1	1	1	-1	1	-1
15	1	-1	1	1	-1	-1	1	1	1	1	-1	1	-1	1	-1	-1	-1	-1	1
16	1	1	-1	1	-1	1	-1	-1	-1	-1	1	1	-1	1	1	-1	-1	1	1
17	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1

Only allows to estimate **primary factors**, not interactions

Preliminary step for further investigation

# Outline

## ① Design of Experiments

## ② Factorial studies

2-level Factorial Studies

ANOVA

Fractional design and Screening

General factorial designs

## ③ Model Investigation

Designs

Exploiting and Reducing Variance

Discussing the Shape of the Model

## ④ Model Estimation

Optimal Designs

## ⑤ Conclusion

# What about having more than two levels?

Before even considering the generation, how would this be analyzed?

- ANOVA still works and interpretation is OK when there are one (1-way ANOVA) or two (2-way ANOVA) factors (with several levels)
- Otherwise, it is a nightmare to analyze and you should decrease either the number of factors or the number of levels

In term of design, you can still go for all combinations

# General Full Factorial Experiments

```
1 d4 <- fac.design(nfactors= 2 ,replications= 3 ,repeat.only= FALSE ,
2   blocks= 1 , randomize= TRUE ,seed= 17366 ,
3   nlevels=c( 3,5 ), factor.names=list(
4   Size=c("S","M","L"),Color=c("R","G","B","M","Y"))); d4
```

```
1   creating full factorial with 15 runs ...
```

	run.no	run.no.std.rp	Size	Color
4	1	6.1	L	G
5	2	10.1	S	M
6	3	12.1	L	M
7	4	2.1	M	R
8	5	11.1	M	M
9	6	14.1	M	Y
10	7	4.1	S	G
11	8	3.1	L	R
12	9	13.1	S	Y
13	10	15.1	L	Y
14	11	5.1	M	G
15	12	1.1	S	R
16	13	8.1	M	B
17	14	7.1	S	B
18	15	9.1	L	B
19	16	1.2	S	R
20	17	8.2	M	B
21	18	15.2	L	Y

# Reducing the size of such designs

You can still sample from it but the outcome is likely to be **not well balanced**

- $\leadsto$  the **estimation** may **not** be that **good** and probably quite biased because of this 😞

```
1 d4[sample(size=5,replace=FALSE,1:nrow(d4)),]
```

```
1      Size Color
2 29      L      G
3 30      L      M
4 41      L      G
5 3       L      M
6 25      S      B
```

That's why you should **try to reduce** as much as possible the number of **factors** and of **levels** if you can



# Outline

## ① Design of Experiments

## ② Factorial studies

2-level Factorial Studies

ANOVA

Fractional design and Screening

General factorial designs

## ③ Model Investigation

Designs

Exploiting and Reducing Variance

Discussing the Shape of the Model

## ④ Model Estimation

Optimal Designs

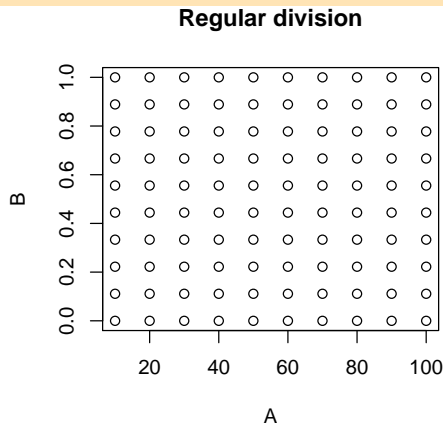
## ⑤ Conclusion

# Without any information about the response

Then we should **not favor a region over an other**

- What about all combinations of a regular division?

```
1 x <- seq(10, 100, length.out = 10)
2 y <- seq(0, 1, length.out = 10)
3 d5_regular <- expand.grid(A = x, B = y)
4 plot(d5_regular, main="Regular division")
```

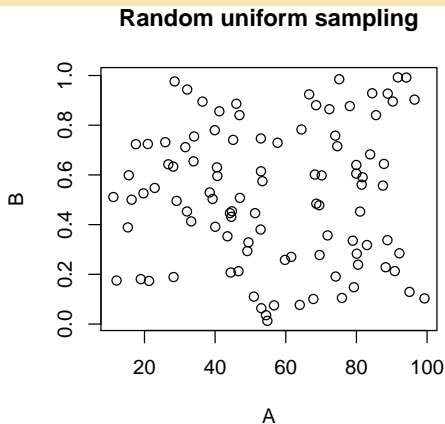


# Can we have a less biased design?

We should **not** favor any particular value

- What about a uniform sampling then?

```
1 set.seed(1);  
2 x <- runif(100,min=10,max=100); y <- runif(100, min=0,max=1)  
3 d5_unif <- data.frame(A = x, B = y)  
4 plot(d5_unif, main="Random uniform sampling")
```




# Can we have a design covering better the whole space?

We do **not** want to **miss any region**

- Space filling designs: **Latin Hyper Square** designs and the **maximin** criteria

```
1 library(DoE.wrapper)
2 d5_maximin <- lhs.design( type= "maximin" , nruns= 100 ,nfactors= 2 ,
3   digits= NULL , seed= 27041 , factor.names=list( A=c(10,100),B=c(0,1) ) )
4 plot(d5_maximin , select = c( "A","B" ), main="LHS design")
```


**LHS design**



# This still reasonably works in higher dimensions

```
1 library(DoE.wrapper); set.seed(42);  
2 d5_HD = lhs.design( type= "maximin" , nruns= 100 ,nfactors= 3 ,  
3   seed= 42 , factor.names=list( A=c(0,1),B=c(0,1),C=c(0,1) ) )  
4 Response5 = 10 + 2*as.numeric(d5_HD$A) + 3*as.numeric(d5_HD$B)*as.numeric(d5_HD$C)  
5 rnorm(nrow(d5_HD),sd=1)  
6 d5_HD <- add.response(d5_HD, Response5, replace=TRUE)  
7 plot(d5_HD , select = c( "A","B","C" ) , main="LHS design")
```

**LHS design**



# What about the analysis?

```
1 summary(lm(Response5 ~ (A + B + C)^2, data = d5_HD))
2
3
4 Call:
5 lm.default(formula = Response5 ~ (A + B + C)^2, data = d5_HD)
6
7 Residuals:
8      Min       1Q   Median       3Q      Max
9 -2.90043 -0.64768  0.00095  0.75471  2.61620
10
11 Coefficients:
12             Estimate Std. Error t value Pr(>|t|)
13 (Intercept)  10.0932     0.5920  17.049  <2e-16 ***
14 A             1.5542     0.9686   1.605   0.1120
15 B             1.1188     0.8904   1.257   0.2121
16 C            -1.4085     0.9283  -1.517   0.1326
17 A:B           -2.3379     1.3228  -1.767   0.0804 .
18 A:C            3.0344     1.2428   2.442   0.0165 *
19 B:C            2.9668     1.2910   2.298   0.0238 *
20 ---
21 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
22 Residual standard error: 1.087 on 93 degrees of freedom
23 Multiple R-squared:  0.451, Adjusted R-squared:  0.4156
24 F-statistic: 12.74 on 6 and 93 DF, p-value: 1.909e-10
```

There is actually too much variability to conclude anything here (look at the  $R^2$ )

We know from the anova that B:C is significant but its Std. Error is still 1.29

We should add another round of 3 times more experiments to halve it

# What happens if we fit a simpler model ?

```
1 summary(lm(Response5 ~ A + B:C, data = d5_HD))
2
3
4 Call:
5 lm.default(formula = Response5 ~ A + B:C, data = d5_HD)
6
7
8 Residuals:
9      Min       1Q   Median       3Q      Max
10 -3.00860 -0.71419 -0.00565  0.74843  2.98579
11
12 Coefficients:
13             Estimate Std. Error t value Pr(>|t|)
14 (Intercept)  10.0054      0.2471  40.489  < 2e-16 ***
15 A              1.8262      0.3920   4.659 1.01e-05 ***
16 B:C           3.0066      0.5247   5.730 1.13e-07 ***
17 ---
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
19
20 Residual standard error: 1.119 on 97 degrees of freedom
21 Multiple R-squared:  0.3938, Adjusted R-squared:  0.3814
22 F-statistic: 31.51 on 2 and 97 DF, p-value: 2.852e-11
```

The Std. Errors decreased but remain quite high

As one could expect, the  $R^2$  has decreased... 😞

# Let's cheat... 😊

```
1 Response5 = 10 + 2*as.numeric(d5_HD$A) + 3*as.numeric(d5_HD$B)*as.numeric(d5_HD$C)
2   rnorm(nrow(d5_HD),sd=.2) # Decreasing variability
3 d5_HD <- add.response(d5_HD, Response5, replace=TRUE)
4 summary(lm(Response5 ~ (A + B + C)^2, data = d5_HD))
```

```
1 Call:
2 lm.default(formula = Response5 ~ (A + B + C)^2, data = d5_HD)
```

```
4 Residuals:
```

```
5      Min       1Q   Median       3Q      Max
6 -0.50030 -0.10491 -0.00945  0.13446  0.47068
```

```
8 Coefficients:
```

```
9             Estimate Std. Error t value Pr(>|t|)
10 (Intercept) 10.06454    0.10992  91.558 < 2e-16 ***
11 A           1.58630    0.17986   8.820 6.41e-14 ***
12 B           0.13805    0.16533   0.835  0.4059
13 C           0.09524    0.17236   0.553  0.5819
14 A:B         0.46421    0.24562   1.890  0.0619 .
15 A:C         0.30745    0.23078   1.332  0.1860
16 B:C         2.33722    0.23972  9.750 6.92e-16 ***
17 ---
```

```
18 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
20 Residual standard error: 0.2019 on 93 degrees of freedom
```

One should actually instead fit the simple model suggested by the previous analysis:

$$y \sim A + B : C$$

You should use **parsimony** both in experiment design and modeling



# Parsimony (1/2)

```
1 summary(lm(Response5 ~ A + B:C, data = d5_HD))
2
3 Call:
4 lm.default(formula = Response5 ~ A + B:C, data = d5_HD)
5
6 Residuals:
7     Min       1Q   Median       3Q      Max
8 -0.56483 -0.11393  0.00626  0.12994  0.46614
9
10 Coefficients:
11             Estimate Std. Error t value Pr(>|t|)
12 (Intercept) 10.05536    0.04609   218.18  <2e-16 ***
13 A             1.94985    0.07311    26.67  <2e-16 ***
14 B:C           2.90476    0.09786    29.68  <2e-16 ***
15 ---
16 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
17
18 Residual standard error: 0.2087 on 97 degrees of freedom
19 Multiple R-squared:  0.95, Adjusted R-squared:  0.949
20 F-statistic: 921.8 on 2 and 97 DF, p-value: < 2.2e-16
```

## Parsimony (2/2)

The principle of **parsimony** is attributed to the 14th century English philosopher **William of Occam**:

*“Given a set of equally good explanations for a given phenomenon, the correct explanation is the simplest explanation”*

## Parsimony (2/2)

The principle of **parsimony** is attributed to the 14th century English philosopher **William of Occam**:

*“Given a set of equally good explanations for a given phenomenon, the correct explanation is the simplest explanation”*

- Models should have **as few parameters as possible**
- Linear models should be preferred to non-linear models
- Models should be **pared down** until they are *minimal adequate*

## Parsimony (2/2)

The principle of **parsimony** is attributed to the 14th century English philosopher **William of Occam**:

*“Given a set of equally good explanations for a given phenomenon, the correct explanation is the simplest explanation”*

- Models should have **as few parameters as possible**
- Linear models should be preferred to non-linear models
- Models should be **pared down** until they are *minimal adequate*

This means, a variable should be retained in the model only if it causes a significant increase in deviance when removed from the current model

*A model should be as simple as possible. But no simpler.*

– A. Einstein

# Outline

## ① Design of Experiments

## ② Factorial studies

2-level Factorial Studies

ANOVA

Fractional design and Screening

General factorial designs

## ③ Model Investigation

Designs

Exploiting and Reducing Variance

Discussing the Shape of the Model

## ④ Model Estimation

Optimal Designs

## ⑤ Conclusion


# Working out a toy example

```
1 x=lhs.design(type= "maximin", nruns=50, nfactors=1, seed=77,  
2     factor.names=list(x=c(0,60)))$x  
3 y=3+x^2/60 + x*rnorm(length(x),sd=.3)  
4 df = data.frame(x=x,y=y)  
5 reg_quad <- lm(data=df,y~x+I(x^2))  
6 summary(reg_quad)
```

```
1 Call:  
2 lm.default(formula = y ~ x + I(x^2), data = df)  
3  
4 Residuals:  
5      Min       1Q   Median       3Q      Max  
6 -21.7802  -4.5247   0.7544   5.1195  20.0284  
7  
8 Coefficients:  
9             Estimate Std. Error t value Pr(>|t|)  
10 (Intercept)  2.124017   4.007473   0.530   0.5986  
11 x           0.143694   0.310362   0.463   0.6455  
12 I(x^2)       0.013169   0.005021   2.623   0.0117 *  
13 ---  
14 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
15  
16 Residual standard error: 9.483 on 47 degrees of freedom  
17 Multiple R-squared:  0.7647, Adjusted R-squared:  0.7547  
18 F-statistic: 76.36 on 2 and 47 DF, p-value: 1.715e-15
```

# We can clearly see where the heteroscedasticity comes from

```
1 xv <- seq(0,60,.5)
2 yv <- predict(reg_quad,list(x=xv,x2=xv^2))
3 ggplot(data=df, aes(x=x,y=y)) + theme_bw() +
4   geom_hline(yintercept=0) + geom_vline(xintercept=0) +
5   geom_point(aes(x=x,y=y)) +
6   geom_line(data=data.frame(x=xv,y=yv),aes(x=x,y=y),color="red")
```



# Adding more points where there is more variability

```
1 x=sqrt(lhs.design( type= "maximin" , nruns= 50 ,nfactors= 1 ,
2     seed= 77 , factor.names=list( x=c(0,60^2) ) )$x)
3 y=3+x^2/60 + x*rnorm(length(x),sd=.3)
4
5 df = data.frame(x=x,y=y)
6 reg_quad <- lm(data=df,y~x+I(x^2))
7 summary(reg_quad)
```

```
1 Call:
2 lm.default(formula = y ~ x + I(x^2), data = df)
```

4 Residuals:

	Min	1Q	Median	3Q	Max
	-27.256	-7.269	1.143	7.702	26.904


8 Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.292996	10.607678	0.028	0.978
x	0.257212	0.626398	0.411	0.683
I(x^2)	0.012031	0.008495	1.416	0.163

14 Residual standard error: 12.02 on 47 degrees of freedom

15 Multiple R-squared: 0.6569, Adjusted R-squared: 0.6423


16 F-statistic: 44.99 on 2 and 47 DF, p-value: 1.209e-11





# Unfortunately, this does not really help

```
1 xv <- seq(0,60,.5)
2 yv <- predict(reg_quad,list(x=xv,x2=xv^2))
3 ggplot(data=df, aes(x=x,y=y)) + theme_bw() +
4   geom_hline(yintercept=0) + geom_vline(xintercept=0) +
5   geom_point(aes(x=x,y=y)) +
6   geom_line(data=data.frame(x=xv,y=yv),aes(x=x,y=y),color="red")
```



The  $R^2$  will never exceed 0.66 because our model fails fully explaining variance

- We should thus rather **replicate** for large values of  $x$  and **average the results**
- The expected value will be the same but the variance will be reduced

# Outline

## ① Design of Experiments

## ② Factorial studies

2-level Factorial Studies

ANOVA

Fractional design and Screening

General factorial designs

## ③ Model Investigation

Designs

Exploiting and Reducing Variance

Discussing the Shape of the Model


## ④ Model Estimation

Optimal Designs

## ⑤ Conclusion

# What if even polynomial models seem inadequate?

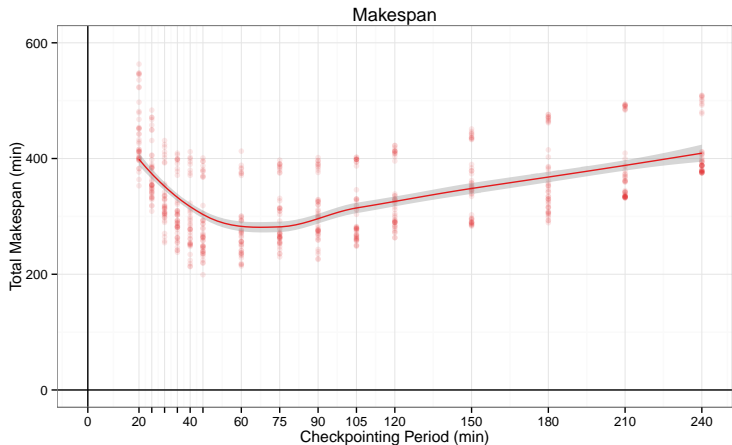
```
1 ggplot(data=df, aes(x=x,y=y)) + theme_bw() +  
2   geom_hline(yintercept=0) + geom_vline(xintercept=0) +  
3   geom_point(aes(x=x,y=y)) +  
4   stat_smooth(method="loess",color="blue",fill="lightblue") +  
5   geom_line(data=data.frame(x=xv,y=yv),aes(x=x,y=y),color="red") +  
6   stat_function(fun=log) # the true function
```



LOcal RegrESSion: builds on linear regression to locally fit a line or a polynomial

This is a *very biased estimator* so use with care

# Discuss about the shape



- "the checkpointing period should be 68 minutes": non-sense, uninteresting 😞
- "optimality region is flat and one should rather overestimate the checkpointing period" 😊

# Outline

## ① Design of Experiments

## ② Factorial studies

2-level Factorial Studies

ANOVA

Fractional design and Screening

General factorial designs

## ③ Model Investigation

Designs

Exploiting and Reducing Variance

Discussing the Shape of the Model

## ④ Model Estimation

Optimal Designs

## ⑤ Conclusion

# D optimality

When estimating model coefficients, it is intuitively better not to spread inputs but rather to use extreme values


- Note: this approach assumes that the model is correct

This intuitive notion can be formalized for linear models (see Hoos):

- Minimize generalized variance of **least squares estimates of model parameters** (determinant of covariance matrix)  
     $\leadsto$  **D-optimal designs**
- Minimize average variance (trace of covariance matrix)  
     $\leadsto$  **A-optimal designs**
- Minimize average of predicted response over experimental region  
     $\leadsto$  **I-optimal designs**


# D-optimal Designs with R

```
1 d7 <- lhs.design(type= "maximin", nruns= 200 , nfactors= 3,  
2   digits=NULL, seed= 20521,  
3   factor.names=list( A=c(0,1),B=c(0,1),C=c(0,1)))  
4 d7.Dopt <- Dopt.design(25, data=d7, formula="~A +B:C", nRepeat= 5,  
5   randomize= TRUE, seed=19583)
```



# D-optimal Designs with R

```
1 d7 <- lhs.design(type= "maximin", nruns= 200 , nfactors= 3,  
2   digits=NULL, seed= 20521,  
3   factor.names=list( A=c(0,1),B=c(0,1),C=c(0,1)))  
4 d7.Dopt <- Dopt.design(25, data=d7, formula="~A +B:C", nRepeat= 5,  
5   randomize= TRUE, seed=19583)
```





# Conclusion

- Designing experiments can be fun! 😊
- Proceed carefully
  - The analysis is not simple but skilled statisticians can help you
  - The **crucial part** is actually the **modeling**, when you identify the factors, the response, and the kind of study
- This lecture only gives an **overview** but may already have **changed your point of view** on how to conduct experiments
- Remind the benefits of the sequential approach:
  - Parsimony
  - Use well-suited DoE and the corresponding analysis
  - Add measurements where there is variability

# Recap on the lecture

- ① Reproducibility is essential
  - literate programming with knitr or org-mode
  - laboratory notebook
- ② Data manipulation and presentation
  - R, ggplot2, plyr, ...
- ③ Introduction to probabilities and statistics
  - A probabilistic model allows you to assess the confidence of your claims
- ④ Linear regression
  - The linear model is quite general
  - This knowledge about the system allows you to improve estimates
- ⑤ Design of Experiments
  - Sequential approach
  - Designs/analysis suited to every study