

Using business workflows to improve control of experiments in distributed systems research

Tomasz Buchert

advised by:

Lucas Nussbaum and Jens Gustedt



Table of Contents

- 1 Introduction
- 2 Contributions
- 3 Details of our approach
- 4 Conclusions

What is a distributed system?



A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.

Leslie Lamport

What is a distributed system? (2)

Examples of well-known distributed systems:

- DNS
- BitTorrent
- Gmail
- HPC applications



What is a distributed system? (3)

Distributed systems are:

- complex
- erroneous
- difficult to control
- nondeterministic

Moreover, they:

- span many networks
- span many geographical locations
- consist of thousands of machines



Research in distributed systems

Can we make better BitTorrent, Gmail, MPI, etc.?

Solution: **experimentation** (like in natural sciences)

The “classic way” is difficult, but other solutions exist:

- simulation
- benchmarking
- emulation

	Real application	Modeled application
Real platform	Standard approach	Benchmarking
Modeled platform	Emulation / Virtualization	Simulation

Research in distributed systems

Can we make better BitTorrent, Gmail, MPI, etc.?

Solution: **experimentation** (like in natural sciences)

The “classic way” is difficult, but other solutions exist:

- simulation
- benchmarking
- emulation

	Real application	Modeled application
Real platform	Standard approach	Benchmarking
Modeled platform	Emulation / Virtualization	Simulation

Experimentation

We all know how frustrating experimenting can be.



That's because experiments in distributed systems are:

- time-consuming
- difficult to do correctly
- complex and incomprehensible
- failure-prone

We have to properly **control** the experiments.

Experimentation tools

Many tools to manage experiments exist:

- Naive way (SSH + Bash + ...)
- Expo
- g5k-campaign
- OMF
- Plush
- ... among many others

They are based on different paradigms.

Our main goal

To improve the research, the experimentation framework has to:

- improve **descriptiveness** of the experiments
- give a **modular** way to build experiments
- handle **unexpected**, but **inevitable** errors
- ensure **scalability** of experiments
- ensure **reproducibility** (or at least **repeatability**)

In the end, we want to

**improve experimentation on testbeds
like Grid'5000 and PlanetLab.**

Bottom-up vs. top-down approach

The majority of tools mentioned use **bottom-up design**.

What about a **top-down** approach?

- 1 Start with high-level description of the experiment.
- 2 Implement low-level details.
- 3 Run the experiment.
- 4 Improve if necessary and reiterate.

There already exists an approach like this.

Business Process Management

Business Process Management is about:

- **understanding** an organization
- modeling its processes as **workflows**
- **executing** and **monitoring** processes
- **improving** organizational **activities**
- redesigning **processes** to make them:
 - cheaper
 - faster
 - less defective



BPM vs. WfM

Business Process Management:

- **not** a technology
- **nothing** to do with computer science

BPM is a **management discipline**.

Workflow Management:

- a set of technologies supporting BPM
- a computer science discipline

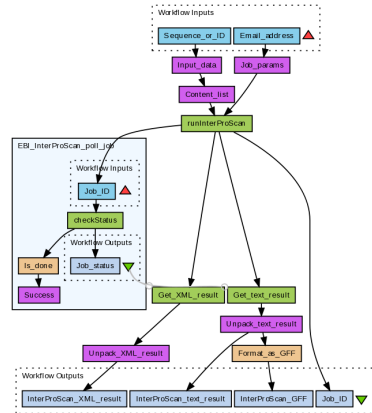
BPM vs. Scientific workflows

Business Process Management:

- workflows describe *control flow*
- arbitrary flow graphs are possible

Scientific workflows:

- *data flows* transforming inputs
- constrained to DAGs
- computing platform is abstracted



The goal of my thesis

**Can BPM improve the experimentation
with distributed systems?
(and if the answer is “yes”, then how?)**

Results

My current work is represented by the following contributions:

- analysis of requirements for an experimentation engine
- development of an approach to control of experiments (and its implementation XPFLOW)
- use of the approach to conduct a large-scale, complicated experiments in a reliable and modular way
- design of a framework to evaluate experiment control tools and evaluation of existing ones

Publications

The following papers are about our approach:

- **A workflow-inspired, modular and robust approach to experiments in distributed systems**
CCGrid 2014, *T. Buchert, L. Nussbaum, J. Gustedt*
- **Orchestration d'expériences à l'aide de processus métier**
CompAS 2013, *T. Buchert, L. Nussbaum*
- **Leveraging business workflows in distributed systems research for the orchestration of reproducible and scalable experiments**
MajecSTIC 2012, *T. Buchert*

The following papers use our approach:

- **Emulation at Very Large Scale with Distem**
CCGrid Scale Challenge 2014, *T. Buchert, E. Jeanvoine, L. Nussbaum*
- **Scalable and Reliable Data Broadcast with Kascade**
HPDIC 2014, *S. Martin, T. Buchert, P. Willemet, O. Richard, E. Jeanvoine, L. Nussbaum*

The remaining papers are:

- **A taxonomy of experiment management tools for distributed systems**
(unpublished yet), *T. Buchert, C. Ruiz, L. Nussbaum, O. Richard*

Publication 1

Leveraging business workflows in distributed systems research for the orchestration of reproducible and scalable experiments

In this article we defined goals and requirements (*desiderata*) for an experiment engine and positioned our approach among existing ones.

We showed that BPM can indeed help.

Publication 1 (cont.)

Design

Descriptiveness
Modularity
Reusability
Maintainability
Support for common
patterns

Execution

Snapshotting
Error handling
Integration with
lower-level tools
Human interaction

Monitoring

Monitoring
Instrumentation
Data analysis

Publication 2

Orchestration d'expériences à l'aide de processus métier

In this article we validated our approach by:

- describing our early implementation
- presenting our language to describe experiments
- testing the new approach with an MPI experiment

Publication 3

A workflow-inspired, modular and robust approach to experiments in distributed systems



In this article we showed a large-scale, cloud-like experiment managed with XPFLOW, by benchmarking Nginx web server with three different testbeds:

- with a local one built with Linux Containers
- with the Grid'5000 testbed
- with a cloud-like testbed (using KVM) hosted on the Grid'5000 testbed

We showed:

- formal description of workflow execution and error handling
- modularity, by easily switching testbeds and composing experiments
- robustness, by handling failures during the experiment

Publication 3 (cont.)



Publication 4

Scalable and Reliable Data Broadcast with Kascade

The article features a comprehensive performance evaluation of a file distribution method.

Nearly all experiments were run with XP_{FLOW}, showing its usefulness.

Publication 4 (cont.)




Fig. 8. The performance of the methods using 10 Gbit/s Ethernet network. No method is able to saturate the available bandwidth.




Fig. 10. The performance of the methods with random ordering of nodes. Kascade relies on a proper ordering of nodes and just like other methods shows low performance in this artificial scenario.

Publication 5

Emulation at Very Large Scale with Distem

The article is a submission to Scale challenge that focuses on very large-scale experiments.

Owing to Distem tool and robustness of XP_{FLOW} control, we achieved an experiment involving 40000 nodes.




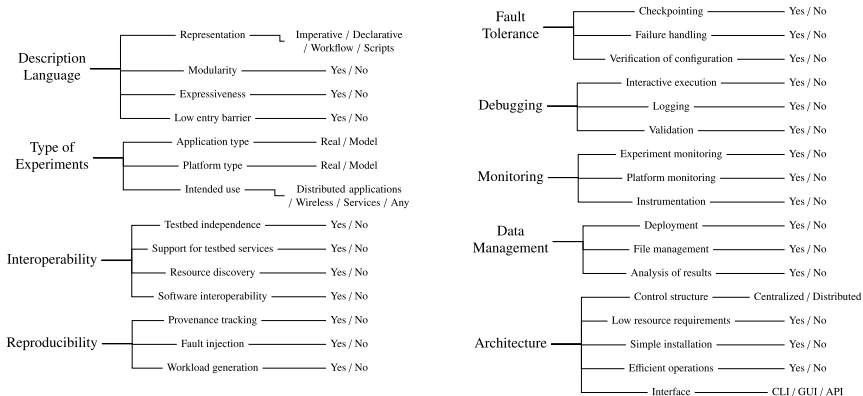
Figure 4. The time required to execute a command using various methods (large number of nodes). Clush is increasingly outperformed by both variations of TakTuk which use a tree overlay to execute commands.

Publication 6


A taxonomy of experiment management tools for distributed systems

The article defines a verbose and useful framework to evaluate experiment control engines like XP_{FLOW} and uses it to evaluate the most prominent ones.

Publication 6 (cont.)




Crash course in workflows

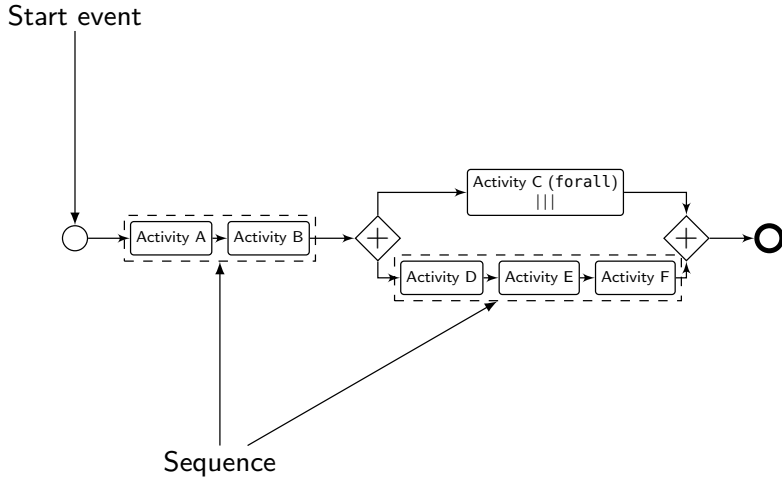


Crash course in workflows

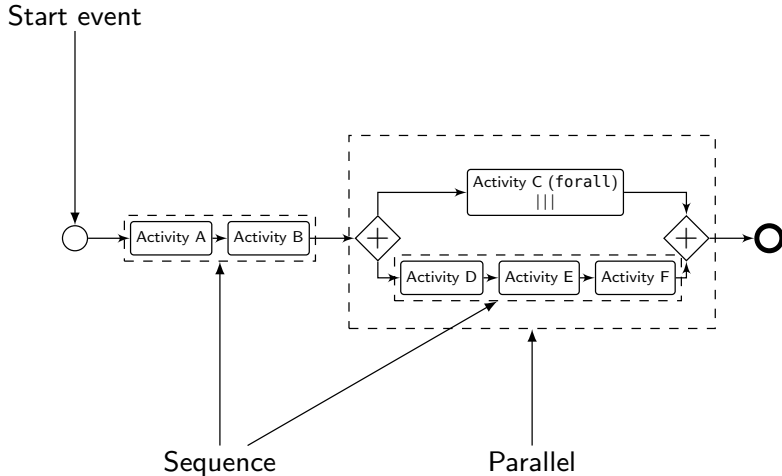
Start event



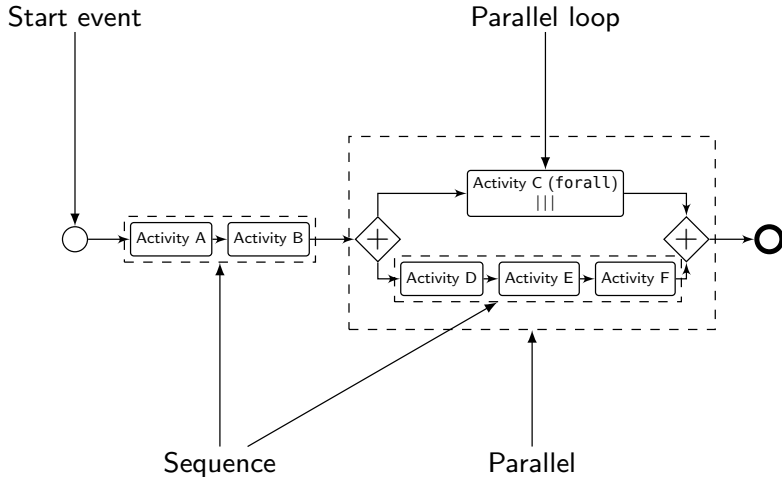
Crash course in workflows



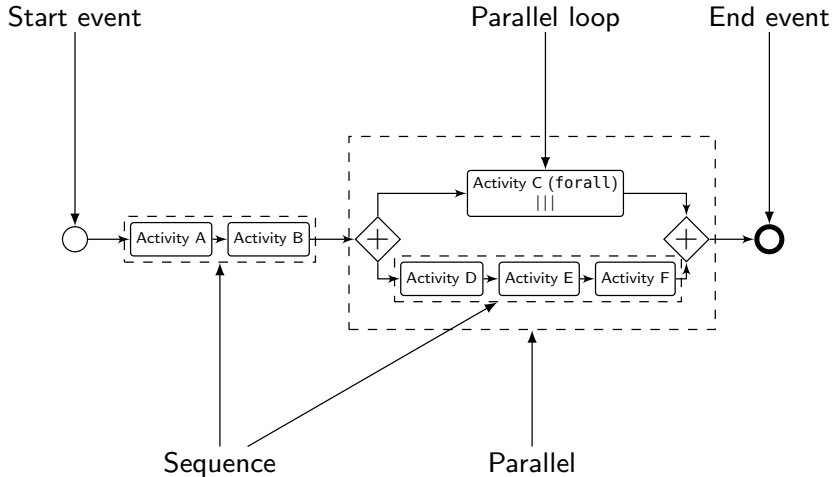
Crash course in workflows



Crash course in workflows




Crash course in workflows



Main concepts

There are 2 main concepts in our approach:


- Processes – high-level description of an experiment:
 - workflows written in a DSL
 - orchestrate other processes and activities
- Activities – low-level building blocks of experiments:
 - do real hard work
 - written in a standard programming language



Main concepts

There are 2 main concepts in our approach:

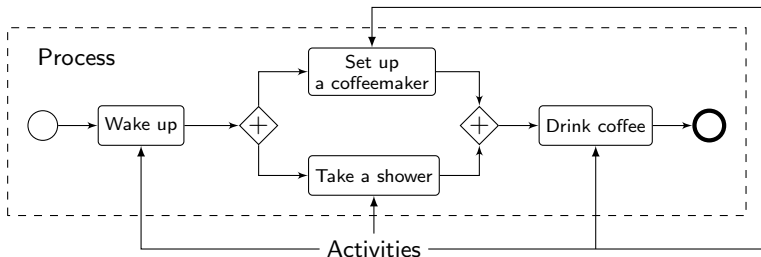
- Processes – high-level description of an experiment:
 - workflows written in a DSL
 - orchestrate other processes and activities
- Activities – low-level building blocks of experiments:
 - do real hard work
 - written in a standard programming language



Main concepts

There are 2 main concepts in our approach:

- Processes – high-level description of an experiment:
 - workflows written in a DSL
 - orchestrate other processes and activities
- Activities – low-level building blocks of experiments:
 - do real hard work
 - written in a standard programming language




Domain-specific language for processes

The DSL for processes features different **workflow patterns**:

- running activities and other processes (run),
- running activities in order or in parallel (sequence, parallel),
- conditional expressions (if, switch)
- running sequential and parallel loops (loop, foreach, forall),
- error handling (try, checkpoint).

Some of them are taken directly from BPM.

DSL example




```

process :workflow do |array|
  run :a
  run :b
  parallel do
    forall array do |x|
      run :c, x
    end
    sequence do
      run :d
      run :e
      run :f
    end
  end
end
end
end

```

DSL example




```

process :workflow do |array|
  run :a
  run :b
  parallel do
    forall array do |x|
      run :c, x
    end
    sequence do
      run :d
      run :e
      run :f
    end
  end
end
end

```

DSL example

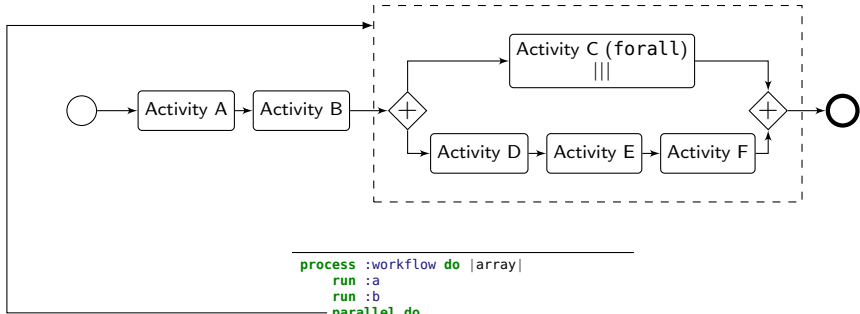


```

process :workflow do |array|
  run :a
  run :b
  parallel do
    forall array do |x|
      run :c, x
    end
    sequence do
      run :d
      run :e
      run :f
    end
  end
end
end

```

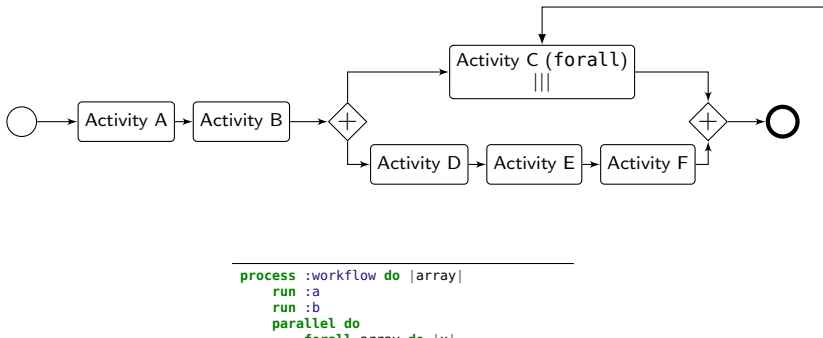

DSL example



```

process :workflow do |array|
  run :a
  run :b
  parallel do
    forall array do |x|
      run :c, x
    end
    sequence do
      run :d
      run :e
      run :f
    end
  end
end
end
  
```

DSL example




```

process :workflow do |array|
  run :a
  run :b
  parallel do
    forall array do |x|
      run :c, x
    end
    sequence do
      run :d
      run :e
      run :f
    end
  end
end
end

```

DSL example

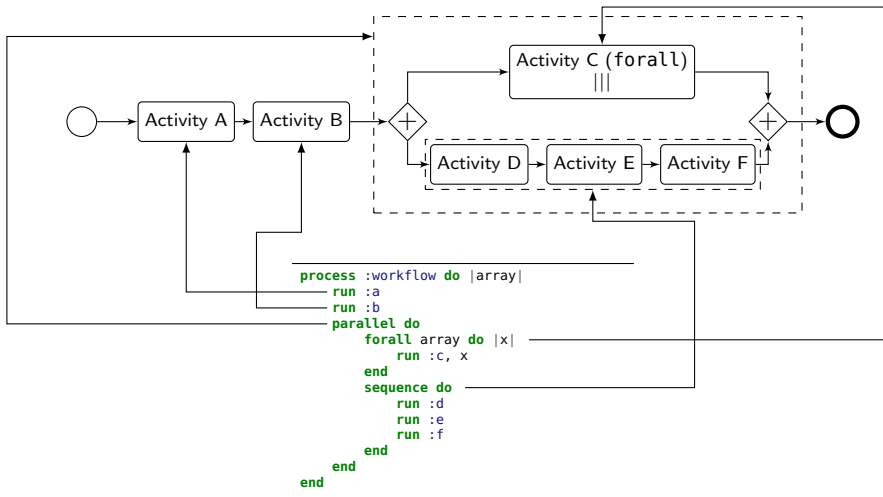


```

process :workflow do |array|
  run :a
  run :b
  parallel do
    forall array do |x|
      run :c, x
    end
    sequence do
      run :d
      run :e
      run :f
    end
  end
end
end
end

```

DSL example



Error handling

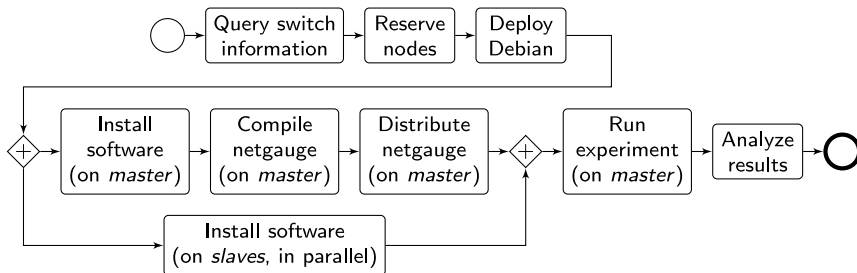
XPFLOW gives some means to cope with failures:

- snapshotting:
 - saves a state of an experiment for future use
 - shortens a development's cycle
- retry policy:
 - retries a failed subprocess execution
 - manages timeouts of operations
 - improves reliability

```
process :snapshotting do
  run :long_deployment
  checkpoint :d
  run :experiment
end
```

```
process :retrying do
  try :retry => 5 do
    run :tricky_activity
  end
end
```

An experiment (once again)



An experiment workflow - DSL representation

```

process :exp do |site, switch|
  s = run g5k.switch, site, switch
  ns = run g5k.nodes, s
  r = run g5k.reserve_nodes,
    :nodes => ns, :time => '2h',
    :site => site, :type => :deploy
  master = (first_of ns)
  rest = (tail_of ns)
  run g5k.deploy,
    r, :env => 'squeeze-x64-nfs'
  checkpoint :deployed
  parallel :retry => true do
    forall rest do |slave|
      run :install_pkgs, slave
    end
    sequence do
      run :install_pkgs, master
      run :build_netgauge, master
      run :dist_netgauge,
        master, rest
    end
  end
  checkpoint :prepared
  output = run :netgauge, master, ns
  checkpoint :finished
  run :analysis, output, switch
end

```

An experiment workflow - DSL representation

```

process :exp do |site, switch|
  s = run g5k.switch, site, switch
  ns = run g5k.nodes, s
  r = run g5k.reserve_nodes,
    :nodes => ns, :time => '2h',
    :site => site, :type => :deploy
  master = (first_of ns)
  rest = (tail_of ns)
  run g5k.deploy,
    r, :env => 'squeeze-x64-nfs'
  checkpoint :deployed
  parallel :retry => true do
    forall rest do |slave|
      run :install_pkgs, slave
    end
    sequence do
      run :install_pkgs, master
      run :build_netgauge, master
      run :dist_netgauge,
        master, rest
    end
  end
end
checkpoint :prepared
output = run :netgauge, master, ns
checkpoint :finished
run :analysis, output, switch
end

```

Activity :install_pkgs

```

activity :install_pkgs do |node|
  log 'Installing packages on ', node
  run 'g5k.bash', node do
    aptget :update
    aptget :upgrade
    aptget :purge, 'mx'
  end
end

```


An experiment workflow - DSL representation

```

process :exp do |site, switch|
  s = run g5k.switch, site, switch
  ns = run g5k.nodes, s
  r = run g5k.reserve_nodes,
    :nodes => ns, :time => '2h',
    :site => site, :type => :deploy
  master = (first_of ns)
  rest = (tail_of ns)
  run g5k.deploy,
    r, :env => 'squeeze-x64-nfs'
  checkpoint :deployed
  parallel :retry => true do
    forall rest do |slave|
      run :install_pkgs, slave
    end
    sequence do
      run :install_pkgs, master
      run :build_netgauge, master
      run :dist_netgauge,
        master, rest
    end
  end
end
checkpoint :prepared
output = run :netgauge, master, ns
checkpoint :finished
run :analysis, output, switch
end

```

Activity :build_netgauge

```

activity :build_netgauge do |master|
  log "Building netgauge on #{master}"
  run 'g5k.copy', NETGAUGE, master, ''
  run 'g5k.bash', master do
    build_tarball NETGAUGE, PATH
  end
  log "Build finished."
end

```

An experiment workflow - DSL representation

```

process :exp do |site, switch|
  s = run g5k.switch, site, switch
  ns = run g5k.nodes, s
  r = run g5k.reserve_nodes,
      :nodes => ns, :time => '2h',
      :site => site, :type => :deploy
  master = (first_of ns)
  rest = (tail_of ns)
  run g5k.deploy,
      r, :env => 'squeeze-x64-nfs'
  checkpoint :deployed
  parallel :retry => true do
    forall rest do |slave|
      run :install_pkgs, slave
    end
    sequence do
      run :install_pkgs, master
      run :build_netgauge, master
      run :dist_netgauge,
          master, rest
    end
  end
end
checkpoint :prepared
output = run :netgauge, master, ns
checkpoint :finished
run :analysis, output, switch
end

```

Activity :dist_netgauge

```

activity :dist_netgauge do |m, s|
  master, slaves = m, s
  run 'g5k.dist_keys', master, slaves
  run 'g5k.bash', master do
    distribute BINARY,
        DEST, 'localhost', slaves
  end
end

```

An experiment workflow - DSL representation

```

process :exp do |site, switch|
  s = run g5k.switch, site, switch
  ns = run g5k.nodes, s
  r = run g5k.reserve_nodes,
      :nodes => ns, :time => '2h',
      :site => site, :type => :deploy
  master = (first_of ns)
  rest = (tail_of ns)
  run g5k.deploy,
      r, :env => 'squeeze-x64-nfs'
  checkpoint :deployed
  parallel :retry => true do
    forall rest do |slave|
      run :install_pkgs, slave
    end
    sequence do
      run :install_pkgs, master
      run :build_netgauge, master
      run :dist_netgauge,
          master, rest
    end
  end
  checkpoint :prepared
  output = run :netgauge, master, ns
  checkpoint :finished
  run :analysis, output, switch
end

```

Activity :netgauge

```

activity :netgauge do |master, nodes|
  log "Running experiment..."
  out = run 'g5k.bash', master do
    cd PATH
    mpirun nodes, "./netgauge"
  end
  log "Experiment done."
end

```

Current activities

Among my current activities are:

- validating the approach with large-scale experiments:
 - the advantages of workflow representation
 - ensuring reliability of experiment execution
 - interoperability with external tools and services
- provenance of data in experiment (related to data analysis)
- provenance of experiment description
- early validation of experiment description
- preparing release of XPFLOW

Provenance

Provenance tries to answer many questions:

- how the data were produced?
- when the data were produced?
- which nodes were involved?

Provenance information can be used to:

- verify the experimental results
- reproduce them
- find problems with the platform (e.g., a reason for outliers)
- cache parts of the experiment

Provenance (cont.)

Provenance in *data flows* (e.g., scientific workflows) is (mostly) a solved problem:

- DAG nodes have a well-defined *parent-child* relation
- data flows are (generally) deterministic
- data is a primary concern of workflows

As far as we know, the provenance collection is nearly non-existing in BPM tools and approaches.

My current work is to introduce provenance to our BPM-based approach.

Conclusions

In this talk I presented my work:

- my research topic and the main goals
- 6 related publications
- overview of the implementation (XPFLOW)
- my current activities

Future work:

- extending the idea to provide reproducibility, provenance, design of experiments, easy data analysis and visualization, validation, monitoring, etc.
- work on lower-level services for efficient and scalable experiments
- release of our software (during the next Grid'5000 school)

Thank you for your attention. Questions?